

# Week 3: Data Tables

Ellen Bledsoe

2024-01-30

## 2-dimensional Data and the tidyverse

### The tidyverse: What is it?

Different programming languages have different syntax (language structure). The **tidyverse** is a package (more accurately, a set of packages) offered in R that all have similar goals and a unified syntax designed to work particularly well with 2-dimensional data (data with rows and column). We call these dataframes in R.

Until now, all of the coding we have done is in the original R language, which is often called “base R.” The syntax in the **tidyverse** is often pretty different from base R. Both are useful, and many people often combine them, which is why we start with base R.

### Wait, what is a package??

Packages are one of the neatest things of working in an open-source environment like R! They contains bits of code (often in the form of functions) that can be reused, making them a core component of reproducible data science. Anyone can develop a package, and there are thousands of them doing all sorts of things.

### Explore the tidyverse

If you want to learn more about the tidyverse, head over to [www.tidyverse.org](http://www.tidyverse.org) and browse the site. Below is a brief summary of *some* of the packages I think you might find the most useful.

- **tidyr**: creating data that is consistent in form/shape
- **dplyr**: creating data that is clean, easily wrangled, and summarized
- **ggplot2**: publication-worthy plots using The Grammar of Graphics
- **tibble**: data frames but better!
- **readr**: fast and friendly ways to read data into R
- **stringr**: easy manipulation of strings (character data)
- **lubridate**: easy manipulation of time and date values

### Using dplyr and readr

#### Download and install

In most scenarios, you will need to download a package from the internet onto your computer before you can use it in RStudio. However, with Posit Cloud, I’ve already done this step for you!

For future reference, though:

- you usually only need to go through this process once until you update R
- we use the function `install.packages()` to download the package

```
# download and install dplyr and readr
# to run the line of code, remove the # in front of the line below and run this chunk
# install.packages("dplyr")
# install.packages("readr")
```

## Load into R

Any time we open R/RStudio and want to use functions from a package, we need to “load” the package. We use the `library()` function to do this.

```
# load the tidyverse (tell RStudio we want to use this package in this session)
library(readr)
library(dplyr)
```

## Set-Up

We are going to download to files into Posit Cloud for us to work with this week. Go ahead and run this code chunk. You should see new CSV files show up in the Files tab.

```
# for the lesson
download.file("https://ndownloader.figshare.com/files/2292172", "surveys.csv")
download.file("https://ndownloader.figshare.com/files/3299474", "plots.csv")
download.file("https://ndownloader.figshare.com/files/3299483", "species.csv")

# for the assignment
download.file("http://www.datacarpentry.org/semester-biology/data/shrub-volume-data.csv", "shrub-volume-assignment.csv")
```

We’ve already talked about some of the benefits of CSV file. They work very nicely in R.

Click on `species.csv` and select View File. If we look at one of these files we can see that it is plain text, so any program can read it. This makes it *interoperable*, which is an important tenant of reproducibility.

The first row is the header row, with different column headers separated by commas. All of the other rows are the data, again with different columns separated by commas. Hence the name “comma separated values.”

## Loading and Viewing the Data

We load these into R using a function from the `readr` package called `read_csv()`.

```
surveys <- read_csv("surveys.csv")

## Rows: 35549 Columns: 9
## -- Column specification -----
## Delimiter: ","
## chr (2): species_id, sex
## dbl (7): record_id, month, day, year, plot_id, hindfoot_length, weight
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
species <- read_csv("species.csv")
```

```
## Rows: 54 Columns: 4
## -- Column specification -----
## Delimiter: ","
## chr (4): species_id, genus, species, taxa
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
plots <- read_csv("plots.csv")
```

```
## Rows: 24 Columns: 2
## -- Column specification -----
## Delimiter: ","
## chr (1): plot_type
## dbl (1): plot_id
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

The tidyverse (in this case, the `read_csv()` function) converts 2D data into something called a tibble! For our intents and purposes, it is basically the same as a data frame (and I'll probably call it a data frame, in reality).

We have three tables here:

- **surveys**: main table, one row for each rodent captured, date on date, location, species ID, sex, and size
- **species**: Latin species names for each species ID and general taxonomic information
- **plots**: information on the experimental manipulations at the site

A few things to note about these tables:

- Good tabular data structure; one table per type of data
- Tables can be linked together to combine information.
- Each row contains a single record (single observation or data point)
- Each column or field contains a single attribute or type of information

We can explore the data frames in the Environment tab or through some functions.

```
str(species)
```

```
## spc_tbl_ [54 x 4] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ species_id: chr [1:54] "AB" "AH" "AS" "BA" ...
## $ genus      : chr [1:54] "Amphispiza" "Ammospermophilus" "Ammodramus" "Baiomys" ...
## $ species    : chr [1:54] "bilineata" "harrisi" "savannarum" "taylori" ...
## $ taxa       : chr [1:54] "Bird" "Rodent" "Bird" "Rodent" ...
## - attr(*, "spec")=
## .. cols(
## ..   species_id = col_character(),
```

```
## .. genus = col_character(),
## .. species = col_character(),
## .. taxa = col_character()
## .. )
## - attr(*, "problems")=<externalptr>
```

```
names(species)
```

```
## [1] "species_id" "genus"      "species"    "taxa"
```

```
head(species)
```

```
## # A tibble: 6 x 4
##   species_id genus      species      taxa
##   <chr>      <chr>      <chr>      <chr>
## 1 AB        Amphispiza  bilineata  Bird
## 2 AH        Ammospermophilus harrisi    Rodent
## 3 AS        Ammodramus   savannarum Bird
## 4 BA        Baiomys     taylori    Rodent
## 5 CB        Campylorhynchus brunneicapillus Bird
## 6 CM        Calamospiza  melanocorys Bird
```

```
glimpse(species) # from dplyr
```

```
## Rows: 54
## Columns: 4
## $ species_id <chr> "AB", "AH", "AS", "BA", "CB", "CM", "CQ", "CS", "CT", "CU", ~
## $ genus      <chr> "Amphispiza", "Ammospermophilus", "Ammodramus", "Baiomys", ~
## $ species    <chr> "bilineata", "harrisi", "savannarum", "taylori", "brunneica~
## $ taxa       <chr> "Bird", "Rodent", "Bird", "Rodent", "Bird", "Bird", "Bird", ~
```

## Intro to dplyr

dplyr (pronounced D-ply-R) is a modern data manipulation library for R that uses the syntax of the tidyverse. It tends to be a bit more intuitive than using base R, especially as tasks become more complicated.

### select()ing columns

Let's use our first function, `select()`. Select allows us to pick out specific columns from our data. You can use names or their position in the data frame.

The first argument in the function is the data frame. Any following arguments are the columns we want to select.

```
# select one column
select(surveys, year)
```

```
## # A tibble: 35,549 x 1
##   year
##   <dbl>
## 1  1977
## 2  1977
## 3  1977
## 4  1977
## 5  1977
## 6  1977
## 7  1977
## 8  1977
## 9  1977
## 10 1977
## # i 35,539 more rows
```

```
# select multiple columns (in any order)
select(surveys, year, month, day)
```

```
## # A tibble: 35,549 x 3
##   year month   day
##   <dbl> <dbl> <dbl>
## 1  1977     7    16
## 2  1977     7    16
## 3  1977     7    16
## 4  1977     7    16
## 5  1977     7    16
## 6  1977     7    16
## 7  1977     7    16
## 8  1977     7    16
## 9  1977     7    16
## 10 1977     7    16
## # i 35,539 more rows
```

```
select(surveys, month, day, year)
```

```
## # A tibble: 35,549 x 3
##   month   day year
##   <dbl> <dbl> <dbl>
## 1     7    16 1977
## 2     7    16 1977
## 3     7    16 1977
## 4     7    16 1977
## 5     7    16 1977
## 6     7    16 1977
## 7     7    16 1977
## 8     7    16 1977
## 9     7    16 1977
## 10    7    16 1977
## # i 35,539 more rows
```

```
# select a consecutive columns
select(surveys, month:year)
```

```
## # A tibble: 35,549 x 3
##   month   day year
##   <dbl> <dbl> <dbl>
## 1     7    16 1977
## 2     7    16 1977
## 3     7    16 1977
## 4     7    16 1977
## 5     7    16 1977
## 6     7    16 1977
## 7     7    16 1977
## 8     7    16 1977
## 9     7    16 1977
## 10    7    16 1977
## # i 35,539 more rows
```

```
# remove columns
select(surveys, -weight)
```

```
## # A tibble: 35,549 x 8
##   record_id month   day year plot_id species_id sex hindfoot_length
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr> <chr> <dbl>
## 1         1     7    16 1977     2 NL      M         32
## 2         2     7    16 1977     3 NL      M         33
## 3         3     7    16 1977     2 DM      F         37
## 4         4     7    16 1977     7 DM      M         36
## 5         5     7    16 1977     3 DM      M         35
## 6         6     7    16 1977     1 PF      M         14
## 7         7     7    16 1977     2 PE      F         NA
## 8         8     7    16 1977     1 DM      M         37
## 9         9     7    16 1977     1 DM      F         34
## 10        10     7    16 1977     6 PF      F         20
## # i 35,539 more rows
```

It is important to remember that the computer interprets everything literally. We need to tell the function the **exact** names of the columns.

## Let's Practice!

Get started with your Assignment. After the set-up, work on Shrub Volume Data Basics a-b.

## Creating new columns with `mutate()`

Sometimes our data doesn't have our data in exactly the format we want. For example, we might want our hindfoot data in cm instead of mm.

The `dplyr` function called `mutate()` lets us create a new column.

The first part of the argument in the `mutate` function (before the `=`) is the name of the new column we want to create (or, sometimes, the name of a column we want to overwrite). After the `=` is what we want the new column to contain.

```
mutate(surveys, hindfoot_length_cm = hindfoot_length / 10)
```

```
## # A tibble: 35,549 x 10
##   record_id month   day  year plot_id species_id sex  hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl>   <dbl> <chr>      <chr>      <dbl>   <dbl>
## 1         1     7    16  1977     2  NL        M          32     NA
## 2         2     7    16  1977     3  NL        M          33     NA
## 3         3     7    16  1977     2  DM        F          37     NA
## 4         4     7    16  1977     7  DM        M          36     NA
## 5         5     7    16  1977     3  DM        M          35     NA
## 6         6     7    16  1977     1  PF        M          14     NA
## 7         7     7    16  1977     2  PE        F          NA     NA
## 8         8     7    16  1977     1  DM        M          37     NA
## 9         9     7    16  1977     1  DM        F          34     NA
## 10        10     7    16  1977     6  PF        F          20     NA
## # i 35,539 more rows
## # i 1 more variable: hindfoot_length_cm <dbl>
```

If we look at the `surveys` object, will it contain the new column?

To store the results of these functions for later, use we need to assign them to a new object or overwrite the existing object.

```
surveys_cm <- mutate(surveys, hindfoot_length_cm = hindfoot_length / 10)
```

## Sorting data with `arrange()`

We can sort the data in the table using the `arrange()` function. Let's sort the `surveys` table by weight.

```
arrange(surveys, weight)
```

```
## # A tibble: 35,549 x 9
##   record_id month   day  year plot_id species_id sex  hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl>   <dbl> <chr>      <chr>      <dbl>   <dbl>
## 1         218     9    13  1977     1  PF        M          13     4
## 2        4052     4     5  1981     3  PF        F          15     4
## 3        4290     4     6  1981     4  PF        <NA>       NA     4
## 4        5346     2    22  1982    21  PF        F          14     4
## 5        7084    11    22  1982     3  PF        F          16     4
## 6        8736    12     8  1983    19  RM        M          17     4
## 7        9790     1    19  1985    16  RM        F          16     4
## 8        9794     1    19  1985    24  RM        M          16     4
## 9        9799     1    19  1985    19  RM        M          16     4
## 10       9823     1    19  1985    23  RM        M          16     4
## # i 35,539 more rows
```

We can see that the rows are now in order from the smallest weight to the largest.

We can reverse the order of the sort by “wrapping” weight in another function: `desc()` for “descending”

```
arrange(surveys, desc(weight))
```

```
## # A tibble: 35,549 x 9
##   record_id month   day  year plot_id species_id sex  hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr>    <chr>          <dbl> <dbl>
## 1     33049    11    17  2001     12  NL      M             33    280
## 2     12871     5    28  1987      2  NL      M             32    278
## 3     15459     1    11  1989      9  NL      M             36    275
## 4       2133    10    25  1979      2  NL      F             33    274
## 5     12729     4    26  1987      2  NL      M             32    270
## 6     13114     7    26  1987      2  NL      M             NA    269
## 7     30175     1     8  2000      2  NL      M             34    265
## 8       4962    11    22  1981     12  NL      F             NA    264
## 9     12602     4     6  1987      2  NL      M             34    260
## 10    13025     7     1  1987      2  NL      M             33    260
## # i 35,539 more rows
```

We can also sort by multiple columns. Perhaps we want to sort first by `plot_id` and then by the date.

```
arrange(surveys, plot_id, year, month, day)
```

```
## # A tibble: 35,549 x 9
##   record_id month   day  year plot_id species_id sex  hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr>    <chr>          <dbl> <dbl>
## 1         6     7    16  1977      1  PF      M             14    NA
## 2         8     7    16  1977      1  DM      M             37    NA
## 3         9     7    16  1977      1  DM      F             34    NA
## 4        78     8    19  1977      1  PF      M             16     9
## 5        80     8    19  1977      1  DS      M             48    NA
## 6       218     9    13  1977      1  PF      M             13     4
## 7       222     9    13  1977      1  DS      M             52    NA
## 8       239     9    13  1977      1  DS      M             48    NA
## 9       263    10    16  1977      1  DM      M             37    40
## 10      270    10    16  1977      1  DM      F             36    38
## # i 35,539 more rows
```

## Let's Practice

Work on Shrub Volume Data Basics c-d.

## filter()ing rows

The `filter()` function allows you filter rows by certain conditions.

We start with the data frame, then we set a condition that has to be met. Let's say we want only rows for the species ID "DS."

To set a condition, we start with the name of the column the want to filter based upon, `species_id`. We then use `==` to set the condition.

Our condition is that we want rows with the value "DS" in the `species_id` column. "DS" here is a string (character data), not a variable or a column name, so we enclose it in quotation marks.



```
filter(surveys, species_id == "DS")
```

```
## # A tibble: 2,504 x 9
##   record_id month   day year plot_id species_id sex hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr>      <chr>      <dbl> <dbl>
## 1         11     7    16  1977     5 DS        F         53    NA
## 2         17     7    16  1977     3 DS        F         48    NA
## 3         20     7    17  1977    11 DS        F         48    NA
## 4         30     7    17  1977    10 DS        F         52    NA
## 5         42     7    18  1977    18 DS        F         46    NA
## 6         58     7    18  1977    12 DS        M         45    NA
## 7         73     8    19  1977     3 DS        F         44    NA
## 8         76     8    19  1977     9 DS        F         47    NA
## 9         80     8    19  1977     1 DS        M         48    NA
## 10        91     8    20  1977    11 DS        F         50    NA
## # i 2,494 more rows
```

Like with vectors, we can have a condition that is “not equal to” using “!=”. Perhaps we want the data for all species except “DS.”

```
filter(surveys, species_id != "DS")
```

```
## # A tibble: 32,282 x 9
##   record_id month   day year plot_id species_id sex hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr>      <chr>      <dbl> <dbl>
## 1         1     7    16  1977     2 NL        M         32    NA
## 2         2     7    16  1977     3 NL        M         33    NA
## 3         3     7    16  1977     2 DM        F         37    NA
## 4         4     7    16  1977     7 DM        M         36    NA
## 5         5     7    16  1977     3 DM        M         35    NA
## 6         6     7    16  1977     1 PF        M         14    NA
## 7         7     7    16  1977     2 PE        F         NA    NA
## 8         8     7    16  1977     1 DM        M         37    NA
## 9         9     7    16  1977     1 DM        F         34    NA
## 10        10     7    16  1977     6 PF        F         20    NA
## # i 32,272 more rows
```

We can also filter on multiple conditions at once.

In computing we combine conditions in two ways: “and” & “or”

Using “and” means that all of the conditions must be true. Do this in `dplyr`, we can add arguments separated by commas or use the `&`

To get the data on species “DS” for the year 1995:

```
# same thing
filter(surveys, species_id == "DS", year > 1995)
```

```
## # A tibble: 22 x 9
##   record_id month   day year plot_id species_id sex hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr>      <chr>      <dbl> <dbl>
## 1    26304     7     9  1997     2 DS        F         50    113
```

```
## 2      26571      7      29 1997      2 DS      F      49      118
## 3      26792      9      27 1997      2 DS      F      49      113
## 4      26980     10      25 1997      2 DS      F      50      108
## 5      27163     11      22 1997      2 DS      F      50      103
## 6      27306     12      28 1997      2 DS      F      51      111
## 7      27426      1      31 1998      2 DS      F      51      122
## 8      27553      3       1 1998      2 DS      M      50      NA
## 9      27588      3       1 1998      1 DS      M      46      NA
## 10     28013      6      27 1998      2 DS      F      51      106
## # i 12 more rows
```

```
filter(surveys, species_id == "DS" & year > 1995)
```

```
## # A tibble: 22 x 9
##   record_id month   day year plot_id species_id sex hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr>    <chr>      <dbl> <dbl>
## 1     26304     7     9 1997     2 DS      F        50    113
## 2     26571     7    29 1997     2 DS      F        49    118
## 3     26792     9    27 1997     2 DS      F        49    113
## 4     26980    10    25 1997     2 DS      F        50    108
## 5     27163    11    22 1997     2 DS      F        50    103
## 6     27306    12    28 1997     2 DS      F        51    111
## 7     27426     1    31 1998     2 DS      F        51    122
## 8     27553     3     1 1998     2 DS      M        50     NA
## 9     27588     3     1 1998     1 DS      M        46     NA
## 10    28013     6    27 1998     2 DS      F        51    106
## # i 12 more rows
```

This approach is mostly useful for building more complex conditions.

When we want rows for which one or more of the conditions are met, we use “or”, which is the | symbol.

```
# compare to the code above
filter(surveys, species_id == "DS" | year > 1995)
```

```
## # A tibble: 14,817 x 9
##   record_id month   day year plot_id species_id sex hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr>    <chr>      <dbl> <dbl>
## 1      11     7    16 1977     5 DS      F        53     NA
## 2      17     7    16 1977     3 DS      F        48     NA
## 3      20     7    17 1977    11 DS      F        48     NA
## 4      30     7    17 1977    10 DS      F        52     NA
## 5      42     7    18 1977    18 DS      F        46     NA
## 6      58     7    18 1977    12 DS      M        45     NA
## 7      73     8    19 1977     3 DS      F        44     NA
## 8      76     8    19 1977     9 DS      F        47     NA
## 9      80     8    19 1977     1 DS      M        48     NA
## 10     91     8    20 1977    11 DS      F        50     NA
## # i 14,807 more rows
```

```
# if we want multiple options from the same column, need to use "or"
filter(surveys, species_id == "DS" | species_id == "DM" | species_id == "DO")
```

```
## # A tibble: 16,127 x 9
##   record_id month   day  year plot_id species_id sex  hindfoot_length weight
##       <dbl> <dbl> <dbl> <dbl>   <dbl> <chr>      <chr>          <dbl>   <dbl>
## 1         3     7    16  1977     2 DM        F            37      NA
## 2         4     7    16  1977     7 DM        M            36      NA
## 3         5     7    16  1977     3 DM        M            35      NA
## 4         8     7    16  1977     1 DM        M            37      NA
## 5         9     7    16  1977     1 DM        F            34      NA
## 6        11     7    16  1977     5 DS        F            53      NA
## 7        12     7    16  1977     7 DM        M            38      NA
## 8        13     7    16  1977     3 DM        M            35      NA
## 9        14     7    16  1977     8 DM        <NA>          NA      NA
## 10       15     7    16  1977     6 DM        F            36      NA
## # i 16,117 more rows
```

## Let's Practice

Do Shrub Volume Data Basics e-g.

## filter()ing null values

One of the common tasks we use `filter` for is removing null values from data. Based on what we learned before, it's natural to think that we do this by using the condition `weight != NA`.

```
filter(surveys, weight != NA)
```

Why didn't that work? Null values like NA are special. Instead, we use a special function, `is.na()`.

```
filter(surveys, is.na(weight))
```

```
## # A tibble: 3,266 x 9
##   record_id month   day  year plot_id species_id sex  hindfoot_length weight
##       <dbl> <dbl> <dbl> <dbl>   <dbl> <chr>      <chr>          <dbl>   <dbl>
## 1         1     7    16  1977     2 NL        M            32      NA
## 2         2     7    16  1977     3 NL        M            33      NA
## 3         3     7    16  1977     2 DM        F            37      NA
## 4         4     7    16  1977     7 DM        M            36      NA
## 5         5     7    16  1977     3 DM        M            35      NA
## 6         6     7    16  1977     1 PF        M            14      NA
## 7         7     7    16  1977     2 PE        F            NA      NA
## 8         8     7    16  1977     1 DM        M            37      NA
## 9         9     7    16  1977     1 DM        F            34      NA
## 10        10     7    16  1977     6 PF        F            20      NA
## # i 3,256 more rows
```

To remove null values we combine this with `!` for "not"

```
filter(surveys, !is.na(weight))
```

```
## # A tibble: 32,283 x 9
```

```
##      record_id month   day  year plot_id species_id sex  hindfoot_length weight
##      <dbl> <dbl> <dbl> <dbl>   <dbl> <chr>      <chr>          <dbl>   <dbl>
##  1         63     8    19  1977     3 DM        M             35     40
##  2         64     8    19  1977     7 DM        M             37     48
##  3         65     8    19  1977     4 DM        F             34     29
##  4         66     8    19  1977     4 DM        F             35     46
##  5         67     8    19  1977     7 DM        M             35     36
##  6         68     8    19  1977     8 DO        F             32     52
##  7         69     8    19  1977     2 PF        M             15      8
##  8         70     8    19  1977     3 OX        F             21     22
##  9         71     8    19  1977     7 DM        F             36     35
## 10         74     8    19  1977     8 PF        M             12      7
## # i 32,273 more rows
```

It is common to combine a null filter with other conditions using “and” For example we might want all of the data on a species that contains weights.

```
filter(surveys, species_id == "DS", !is.na(weight))
```

```
## # A tibble: 2,344 x 9
##      record_id month   day  year plot_id species_id sex  hindfoot_length weight
##      <dbl> <dbl> <dbl> <dbl>   <dbl> <chr>      <chr>          <dbl>   <dbl>
##  1         357    11    12  1977     9 DS        F             50    117
##  2         362    11    12  1977     1 DS        F             51    121
##  3         367    11    12  1977    20 DS        M             51    115
##  4         377    11    12  1977     9 DS        F             48    120
##  5         381    11    13  1977    17 DS        F             48    118
##  6         383    11    13  1977    11 DS        F             52    126
##  7         385    11    13  1977    17 DS        M             50    132
##  8         389    11    13  1977    14 DS        F             NA    113
##  9         392    11    13  1977    11 DS        F             53    122
## 10         394    11    13  1977     4 DS        F             48    107
## # i 2,334 more rows
```

## Let's Practice

Work on the last of Exercise 1, Shrub Volume.