

Week 7: Character Strings

Ellen Bledsoe

Working with Strings

Working with character strings (text data) can be particularly challenging in R. Thankfully, we have the **stringr** package, also part of the **tidyverse**, to help!

All functions in the **stringr** package start with **str_**. There are *many* helpful functions in the **stringr** package. We'll only cover a handful here, but if you're looking to accomplish something with a string and aren't sure how to approach it, the **stringr** package almost certainly can get it done.

stringr Functions

We'll cover five **stringr** functions in this lesson. Each **stringr** function will take the object (or vector or column...) of character strings as the first argument. From there, the subsequent arguments and the outputs of the functions differ.

Paying attention to the *output* of a **stringr** function is particularly useful in determining when and how to use that function, especially in **filter()** or **mutate()** functions.

Function	Use	Output
str_length()	count the number of characters in the string	numeric vector
str_detect()	determine if pattern is found within string	logical (T/F) vector
str_extract()	return portion of each string that matches the pattern	character vector
str_remove()	remove portion of the string that matches the pattern	character vector
str_replace()	replace portion of string that matches the pattern with something else	character vector

Let's load the **tidyverse** and get started.

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.2.1      v readr      2.2.0
v forcats    1.0.1      v stringr    1.6.0
v ggplot2    4.0.3      v tibble     3.3.1
v lubridate  1.9.5      v tidyr      1.3.2
v purrr      1.2.2
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

We will use some real data from my postdoc lab that I was tasked with cleaning.

It's a long story, but I ended up in an aquatic biogeochemistry lab for my postdoc. The lab had 4+ years of data from dugouts, which are small, human-made water reservoirs that are common across the prairie states and provinces.

Unfortunately, the data were collected by many different people without a standard data entry protocol, so there was extensive data cleaning that needed to happen to bring all of the datasets together across the years.

Many of the tools we are covering in this class (especially the last few weeks with joins, pivots, and strings) were integral to getting those datasets wrangled. Here is just one example of the type of dataset I was working with.

```
url <- "https://raw.githubusercontent.com/bleds22e/FAST_lab_training/master/merging_masters/
carbon <- read_csv(url) %>% rename(SampleID = `Sample ID`)
```

```
Rows: 48 Columns: 3
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
chr (1): Sample ID
```

```
dbl (2): TIC (PPM as mg/L C), TOC (PPM as mg/L C)
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
carbon
```

```
# A tibble: 48 x 3
  SampleID      `TIC (PPM as mg/L C)` `TOC (PPM as mg/L C)`
  <chr>          <dbl>          <dbl>
1 AMC3/JULY13      51.7            33.3
2 BMC3/AUG13      58.3            37.2
3 BMC2/AUG13      89.7            47.6
4 E56A-DEEP/AUG17 259.            57.2
5 ACB2/JUNE22     50.0            22.0
6 E14A/AUG17     107.            56.8
7 CFH/AUG6        22.9            18.8
8 ALA2/JULY22     127.            100.
9 AMC1/JUNE26      48.9            31.8
10 CLH/AUG4        43.4            32.1
# i 38 more rows
```

Strings in Vectors

For simplicity's sake, let's start working with strings first in vectors rather than a data frame.

We want to pull out the first column (the `Sample ID` column) from the `carbon` data frame as a vector; we can do that using the `$` from base R.

```
sampleID <- carbon$SampleID
sampleID
```

```
[1] "AMC3/JULY13"      "BMC3/AUG13"      "BMC2/AUG13"      "E56A-DEEP/AUG17"
[5] "ACB2/JUNE22"      "E14A/AUG17"      "CFH/AUG6"        "ALA2/JULY22"
[9] "AMC1/JUNE26"      "CLH/AUG4"        "A14B/JUNE5"      "BLS/JULY7"
[13] "ALH/JUNE4"        "ALA1/JULY2"      "A56C/JUNE2"      "B14B/JUNE29"
[17] "AFH/MAY29"        "AMC2/JULY13"     "BLA1/AUG12"      "BT1/AUG10"
[21] "C14B/JULY30"      "BLH/JULY7"       "CLS/AUG4"        "ACB1/JUNE22"
[25] "B56A/JUNE30"      "B14A/JUNE29"     "ALHM2/JUNE23"    "AT1/JULY14"
[29] "ALHM1/JUNE23"     "A56A/JUNE2"      "ALS/JUNE4"       "A14A/JUNE5"
[33] "BFH/JULY1"        "BLHM1/AUG11"     "E56A/AUG17"      "BLB1/AUG7"
[37] "D14A/JULY30"      "B56C/JULY29"     "BMC1/AUG13"      "D56A/JULY29"
[41] "E14A-DEEP/AUG17" "AT2/JULY14"      "BT2/AUG10"       "BLHM2/AUG11"
[45] "C14A/JULY16"      "BLA2/AUG12"      "BCB2/AUG7"       "C56A/JULY15"
```

Now, let's demonstrate how each of the five `stringr` functions work. I've listed the *required* arguments in the parentheses for each function.

`str_length(string)`: count the number of characters in the string

```
str_length(sampleID)
```

```
[1] 11 10 10 15 11 10  8 11 11  8 10  9  9 10 10 11  9 11 10  9 11  9  8 11 11
[26] 11 12 10 12 10  9 10  9 11 10  9 11 11 10 11 15 10  9 11 11 10  9 11
```

`str_detect(string, pattern)`: determine if pattern is found within string

```
str_detect(sampleID, "DEEP")
```

```
[1] FALSE FALSE FALSE  TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[13] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[25] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[37] FALSE FALSE FALSE FALSE  TRUE FALSE FALSE FALSE FALSE FALSE FALSE
```

`str_extract(string, pattern)`: return portion of each string that matches the pattern

```
# 3 characters after the /
str_extract(sampleID, "/(...)")
```

```
[1] "/JUL" "/AUG" "/AUG" "/AUG" "/JUN" "/AUG" "/AUG" "/JUL" "/JUN" "/AUG"
[11] "/JUN" "/JUL" "/JUN" "/JUL" "/JUN" "/JUN" "/MAY" "/JUL" "/AUG" "/AUG"
[21] "/JUL" "/JUL" "/AUG" "/JUN" "/JUN" "/JUN" "/JUN" "/JUL" "/JUN" "/JUN"
[31] "/JUN" "/JUN" "/JUL" "/AUG" "/AUG" "/AUG" "/JUL" "/JUL" "/AUG" "/JUL"
[41] "/AUG" "/JUL" "/AUG" "/AUG" "/JUL" "/AUG" "/AUG" "/JUL"
```

```
months <- str_extract(sampleID, "/(...)")
```

`str_remove(string, pattern)`: remove portion of the string that matches the pattern

```
str_remove(months, "/")
```

```
[1] "JUL" "AUG" "AUG" "AUG" "JUN" "AUG" "AUG" "JUL" "JUN" "AUG" "JUN" "JUL"
[13] "JUN" "JUL" "JUN" "JUN" "MAY" "JUL" "AUG" "AUG" "JUL" "JUL" "AUG" "JUN"
[25] "JUN" "JUN" "JUN" "JUL" "JUN" "JUN" "JUN" "JUN" "JUL" "AUG" "AUG" "AUG"
[37] "JUL" "JUL" "AUG" "JUL" "AUG" "JUL" "AUG" "AUG" "JUL" "AUG" "AUG" "JUL"
```

`str_replace(string, pattern, replacement)`: replace portion of string that matches the pattern with something else

```
# we can also replace values (with something new or with a blank)
str_replace(months, "/", "")
```

```
[1] "JUL" "AUG" "AUG" "AUG" "JUN" "AUG" "AUG" "JUL" "JUN" "AUG" "JUN" "JUL"
[13] "JUN" "JUL" "JUN" "JUN" "MAY" "JUL" "AUG" "AUG" "JUL" "JUL" "AUG" "JUN"
[25] "JUN" "JUN" "JUN" "JUL" "JUN" "JUN" "JUN" "JUN" "JUL" "AUG" "AUG" "AUG"
[37] "JUL" "JUL" "AUG" "JUL" "AUG" "JUL" "AUG" "AUG" "JUL" "AUG" "AUG" "JUL"
```

Using stringr in Data Frames

Like we saw with the `lubridate` functions, you will often want to use `stringr` functions within other `tidyverse` functions, such as `filter` or `mutate`.

Let's run through a few examples using the full `carbon` dataframe.

Perhaps we want only samples that were collected in August. We can use the `str_detect` function to set our condition in the `filter` function. This works because the output of `str_detect` is a logical vector (T/F), as with other conditional statements or `is.na()`.

```
carbon %>%
  filter(str_detect(SampleID, "AUG"))
```

```
# A tibble: 18 x 3
  SampleID      `TIC (PPM as mg/L C)` `TOC (PPM as mg/L C)`
  <chr>          <dbl>          <dbl>
1 BMC3/AUG13      58.3            37.2
2 BMC2/AUG13      89.7            47.6
3 E56A-DEEP/AUG17 259.            57.2
4 E14A/AUG17     107.            56.8
5 CFH/AUG6        22.9            18.8
6 CLH/AUG4        43.4            32.1
7 BLA1/AUG12      54.5            71.1
8 BT1/AUG10       71.0            29.9
9 CLS/AUG4        36.4            21.6
10 BLHM1/AUG11    43.9            25.3
11 E56A/AUG17     60.6            29.7
12 BLB1/AUG7      29.4            18.4
13 BMC1/AUG13     49.8            39.2
14 E14A-DEEP/AUG17 111.            55.9
```

15	BT2/AUG10	34.9	15.6
16	BLHM2/AUG11	71.8	37.8
17	BLA2/AUG12	98.1	105.
18	BCB2/AUG7	53.8	24.9

Perhaps we want to filter for sample names over a certain length (for some reason). We can do that as well, though we need to structure our condition a little differently this time, because the output of `str_length` is not a logical vector.

```
carbon %>%
  filter(str_length(SampleID) > 13)
```

```
# A tibble: 2 x 3
  SampleID      `TIC (PPM as mg/L C)` `TOC (PPM as mg/L C)`
  <chr>          <dbl>          <dbl>
1 E56A-DEEP/AUG17      259.          57.2
2 E14A-DEEP/AUG17      111.          55.9
```

Alternatively, perhaps we want to create a column with the month the sample was collected. We can use the `str_extract` and `str_remove` columns in a mutate function.

```
carbon %>%
  mutate(Month = str_extract(SampleID, "/(\\.\\.\\.)"),
         Month = str_remove(Month, "/"))
```

```
# A tibble: 48 x 4
  SampleID      `TIC (PPM as mg/L C)` `TOC (PPM as mg/L C)` Month
  <chr>          <dbl>          <dbl> <chr>
1 AMC3/JULY13      51.7          33.3 JUL
2 BMC3/AUG13       58.3          37.2 AUG
3 BMC2/AUG13       89.7          47.6 AUG
4 E56A-DEEP/AUG17  259.          57.2 AUG
5 ACB2/JUNE22      50.0          22.0 JUN
6 E14A/AUG17      107.          56.8 AUG
7 CFH/AUG6        22.9          18.8 AUG
8 ALA2/JULY22     127.          100. JUL
9 AMC1/JUNE26      48.9          31.8 JUN
10 CLH/AUG4        43.4          32.1 AUG
# i 38 more rows
```

Regular Expressions

While being able to match specific strings is helpful, often we have more complicated requirements, such as counting all the numbers from a string, removing the first 3 characters of a string, or extracting all of the values after a certain symbol.

When we need to perform more complicated tasks using strings, we can turn to something called “regular expressions,” or “regex” for short. Regular expressions uses characters and special symbols to define certain search patterns in concise ways.

I’m not going to go deep into “regex,” but you should know that they exist in case you need to use them in the future.

As one example, let’s say I wanted to pull out all of the characters after the / in the `sampleID` vector, since they represent dates. We could use the regular expression `"(?<=/).*" to do so.`

```
str_extract(sampleID, "(?<=/).*")
```

```
[1] "JULY13" "AUG13" "AUG13" "AUG17" "JUNE22" "AUG17" "AUG6" "JULY22"
[9] "JUNE26" "AUG4" "JUNE5" "JULY7" "JUNE4" "JULY2" "JUNE2" "JUNE29"
[17] "MAY29" "JULY13" "AUG12" "AUG10" "JULY30" "JULY7" "AUG4" "JUNE22"
[25] "JUNE30" "JUNE29" "JUNE23" "JULY14" "JUNE23" "JUNE2" "JUNE4" "JUNE5"
[33] "JULY1" "AUG11" "AUG17" "AUG7" "JULY30" "JULY29" "AUG13" "JULY29"
[41] "AUG17" "JULY14" "AUG10" "AUG11" "JULY16" "AUG12" "AUG7" "JULY15"
```

Helpful Resources

While memorizing regular expressions is wildly daunting, there are thankfully numerous resources that we can use to help us out.

[Here](#) is a website where you can build and test regex.

Honestly, though, I use Gen AI to build my regex!

To build the `"(?<=/).*" expression from above, I asked ChatGPT to “use regex and str_extract to extract everything after a /, not including the /”, and it produced exactly what I needed.`