

Week 14: Dates and Times

Ellen Bledsoe

Working with Dates and Times

This week, we are going to learn a bit about how to work with some of the more troublesome data types in R: dates and character strings.

I've divided the lessons into 2 parts this week. Today, we will work with the `lubridate` package to learn about working with dates and times. The first few questions in the assignment will do the same.

Set-Up

As is becoming standard for us now, let's start by loading in the `tidyverse`.

```
library(tidyverse)
```

We are going to use some stream data from the Santa Cruz River here in Tucson. The data are recorded every 15 minutes; these data are from 2023.

When we take a look at the data, we can see a number of rows that we need to skip when we read in the data...27 to be exact.

```
water <- read_tsv("santa_cruz.txt", skip = 27, col_names = TRUE)

water

# A tibble: 34,704 x 6
  agency_cd site_no   datetime          tz_cd `307121_00060` `307121_00060_cd`
  <chr>     <chr>    <chr>           <chr> <chr>      <chr>
1 5s        15s       20d              6s    14n       10s
2 USGS     09482495  2023-01-01 00:00 MST   0.00      A
3 USGS     09482495  2023-01-01 00:15 MST   0.00      A
```

```

4 USGS      09482495 2023-01-01 00:30 MST  0.00      A
5 USGS      09482495 2023-01-01 00:45 MST  0.00      A
6 USGS      09482495 2023-01-01 01:00 MST  0.00      A
7 USGS      09482495 2023-01-01 01:15 MST  0.00      A
8 USGS      09482495 2023-01-01 01:30 MST  0.00      A
9 USGS      09482495 2023-01-01 01:45 MST  0.00      A
10 USGS     09482495 2023-01-01 02:00 MST  0.00     A
# i 34,694 more rows

```

Ok, we are close but not quite there. The first row of data is metadata rather than data values.

To remove that, we are going to use a new function called `slice()`, from the `dplyr` package. It allows us to choose (or remove, in our case) rows from a dataframe based on position. We want to remove the first row.

```

water <- water %>%
  slice(-1)

water

# A tibble: 34,703 x 6
  agency_cd site_no   datetime          tz_cd `307121_00060` `307121_00060_cd`
  <chr>     <chr>    <chr>           <chr> <chr>       <chr>
1 USGS      09482495 2023-01-01 00:00 MST  0.00      A
2 USGS      09482495 2023-01-01 00:15 MST  0.00      A
3 USGS      09482495 2023-01-01 00:30 MST  0.00      A
4 USGS      09482495 2023-01-01 00:45 MST  0.00      A
5 USGS      09482495 2023-01-01 01:00 MST  0.00      A
6 USGS      09482495 2023-01-01 01:15 MST  0.00      A
7 USGS      09482495 2023-01-01 01:30 MST  0.00      A
8 USGS      09482495 2023-01-01 01:45 MST  0.00      A
9 USGS      09482495 2023-01-01 02:00 MST  0.00      A
10 USGS     09482495 2023-01-01 02:15 MST  0.00     A
# i 34,693 more rows

```

That's better! Let's also rename the last 2 columns. Those numbers are going to be pretty annoying to deal with down the road.

The stream discharge column should be numeric but got read in as character for some reason, so we need to convert it to a numeric data class.

```

water <- water %>%
  rename(dischARGE_ft3_s = `307121_00060`,
         qualification_code = `307121_00060_cd`) %>%
  mutate(dischARGE_ft3_s = as.double(dischARGE_ft3_s))
water

# A tibble: 34,703 x 6
  agency_cd site_no   datetimE       tz_cd  dischARGE_ft3_s qualification_code
  <chr>      <chr>     <chr>       <chr>      <dbl> <chr>
1 USGS      09482495 2023-01-01 00:00 MST          0 A
2 USGS      09482495 2023-01-01 00:15 MST          0 A
3 USGS      09482495 2023-01-01 00:30 MST          0 A
4 USGS      09482495 2023-01-01 00:45 MST          0 A
5 USGS      09482495 2023-01-01 01:00 MST          0 A
6 USGS      09482495 2023-01-01 01:15 MST          0 A
7 USGS      09482495 2023-01-01 01:30 MST          0 A
8 USGS      09482495 2023-01-01 01:45 MST          0 A
9 USGS      09482495 2023-01-01 02:00 MST          0 A
10 USGS     09482495 2023-01-01 02:15 MST          0 A
# i 34,693 more rows

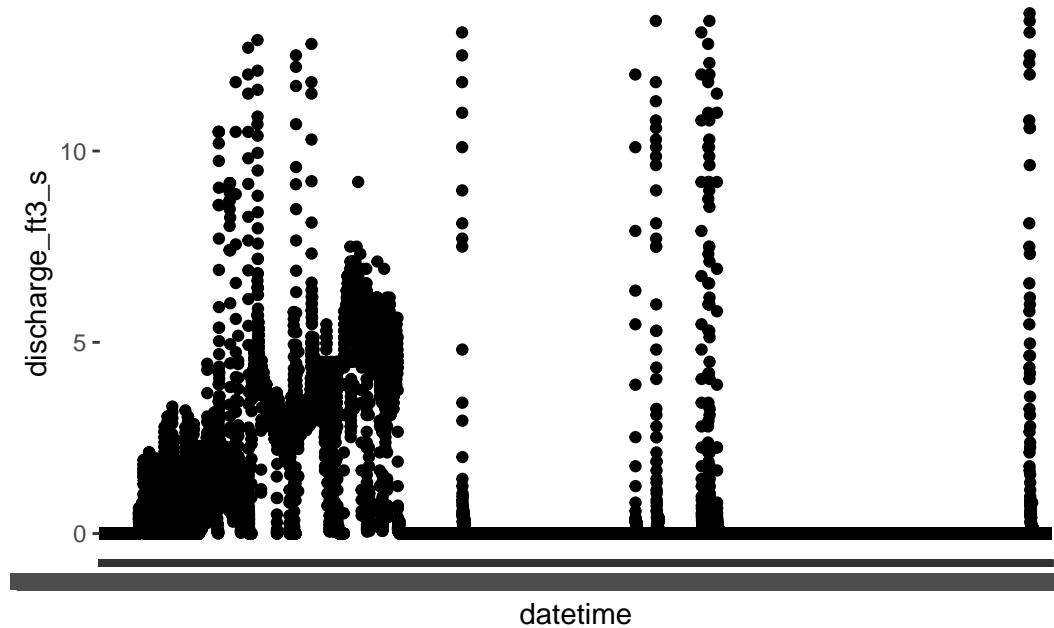
```

Let's try to plot the amount of water flowing through the stream through time.

```

ggplot(water, aes(datetimE, dischARGE_ft3_s)) +
  geom_point()

```



Well, that doesn't look great...

As it turns out, the `datetime` column is also currently character data. In order to use this column efficiently, we will want to convert it to a date (or datetime) data class.

Dates and Times with lubridate

Dates and times can be particularly challenging to work with, but thankfully the `lubridate` package makes it a bit easier.

Referencing the Current Date and Time

First, let's start with functions that return the current date and/or time. These are helpful if you want to "stamp" with the date/time something happens (e.g., when exactly did you knit that PDF file).

```
# date  
today()
```

```
[1] "2025-12-01"
```

```
# date and time  
now()
```

```
[1] "2025-12-01 13:50:11 MST"
```

Aside from these functions, we will mostly be using the `lubridate` functions within a `mutate` function to modify entire columns at once.

Making a `datetime` Column

The most generic function in `lubridate` is `as_date` or the variant, `as_datetime`. These convert non-dates into a date data class.

We need to specify the following arguments:

- the data to be converted to a datetime
- the format of the datetime data
- optionally, we can specify the time zone

The way we specify the format of the datetime is a bit unusual. We use “conversion specifications,” which are introduced by a % and are typically represented by a letter. There are lots of options, many of which are in the image below:

Symbol	Meaning	Example
%a	Abbreviated weekday name	Tue
%A	Full weekday name	Tuesday
%b	Abbreviated month name	Apr
%B	Full month name	April
%C	Century: the integer part of the year divided by 100	20
%d	Day of the month	09
%H	Hours as decimal number (00–23)	13
%I	Hours as decimal number (01–12)	1
%m	Month as number (01–12)	04
%M	Minute as number (00–59)	12
%p	AM/PM indicator for 12-hour time (%I)	PM
%S	Second as integer (00–61)	12
%u	Weekday as a decimal number (1–7, Monday is 1)	2
%w	Weekday as decimal number (0–6, Sunday is 0)	2
%y	2-Digit Year (00–99)	19
%Y	4-Digit Year	2019

Figure 1: Table from: <https://stat.ethz.ch/R-manual/R-devel/library/base/html/strftime.html>

The characters that are used to separate the date and time data are interpreted literally. What does that mean? If the date is separated by -, the conversion specifications should be separated by -, too.

For example, if our dates were formatted as 02/29/24, the corresponding conversion format would be %m/%d/%y.

Let's convert our datetime column to a datetime data class.

```
water %>%
  mutate(datetime = as_datetime(datetime, format = "%Y-%m-%d %H:%M", tz = "MST"))

# A tibble: 34,703 x 6
  agency_cd site_no   datetime           tz_cd discharge_ft3_s
  <chr>      <chr>    <dttm>          <chr>            <dbl>
1 USGS       09482495 2023-01-01 00:00:00 MST              0
2 USGS       09482495 2023-01-01 00:15:00 MST              0
3 USGS       09482495 2023-01-01 00:30:00 MST              0
4 USGS       09482495 2023-01-01 00:45:00 MST              0
5 USGS       09482495 2023-01-01 01:00:00 MST              0
```

```

6 USGS      09482495 2023-01-01 01:15:00 MST      0
7 USGS      09482495 2023-01-01 01:30:00 MST      0
8 USGS      09482495 2023-01-01 01:45:00 MST      0
9 USGS      09482495 2023-01-01 02:00:00 MST      0
10 USGS     09482495 2023-01-01 02:15:00 MST      0
# i 34,693 more rows
# i 1 more variable: qualification_code <chr>

```

We can now see that the `datetime` column is something new, called `POSIXct`.

You don't need to know too much about what that means beyond that this is a common way datetimes are stored in R.

Behind the scenes, the date/time is being stored as the number of seconds since Jan 1, 1970 at 00:00:00 UTC, allowing R to perform mathematical calculations with dates and times, which I'll demonstrate later in the lesson.

If you'd like to learn a bit more about working with `POSIXct` (or other variants), I think [this lesson](#) from NEON does a pretty nice job.

Functions Specifying Date/Time Structure

Personally, any time I try to remember the conversion specifications, I draw a blank. Are they uppercase or lowercase? Does the letter match the first letter of the data it represents? (Spoiler alert: not always...)

`lubridate` provides functions with nearly every iteration and combination of year, month, day, hour, minutes, and seconds.

When we use these, we don't need to specify the exact format or the separators, which is particularly helpful if you have dates with different separators in the same column.

As we know from the previous code, our `datetime` column has year, month, day, hours and minutes.

```

water <- water %>%
  mutate(datetime = ymd_hm(datetime))
water

```

```

# A tibble: 34,703 x 6
  agency_cd site_no   datetime           tz_cd discharge_ft3_s
  <chr>      <chr>    <dttm>          <chr>            <dbl>
1 USGS       09482495 2023-01-01 00:00:00 MST             0
2 USGS       09482495 2023-01-01 00:15:00 MST             0

```

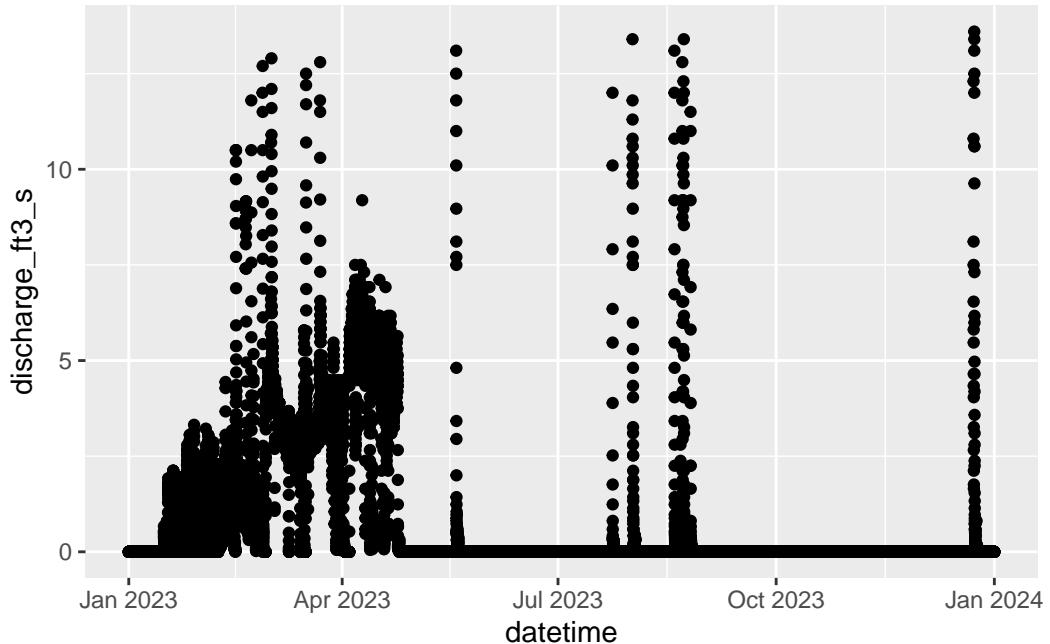
```

3 USGS 09482495 2023-01-01 00:30:00 MST 0
4 USGS 09482495 2023-01-01 00:45:00 MST 0
5 USGS 09482495 2023-01-01 01:00:00 MST 0
6 USGS 09482495 2023-01-01 01:15:00 MST 0
7 USGS 09482495 2023-01-01 01:30:00 MST 0
8 USGS 09482495 2023-01-01 01:45:00 MST 0
9 USGS 09482495 2023-01-01 02:00:00 MST 0
10 USGS 09482495 2023-01-01 02:15:00 MST 0
# i 34,693 more rows
# i 1 more variable: qualification_code <chr>

```

Now, when we plot the trend of water discharge through 2023, we can see that ggplot knows how to deal with the values more efficiently.

```
ggplot(water, aes(datetime, discharge_ft3_s)) +
  geom_point()
```



Extracting Specific Components

Sometimes, we don't want all of the data from the datetime or we want to `filter`, `group_by`, `summarize`, etc. by a part of the datetime.

For example, perhaps we want to know the average water discharge per month. We would first want to create a new column with the month for each observation and then `group_by` and `summarize`.

There are numerous functions we can use to extract specific components from the datetime.

We can use the `date()` function to pull out the date from the datetime. As a note, it is surprisingly challenging to create a time column...

```
water %>%
  mutate(Date = date(datetime))

# A tibble: 34,703 x 7
  agency_cd site_no   datetime           tz_cd discharge_ft3_s
  <chr>      <chr>    <dttm>          <chr>      <dbl>
1 USGS       09482495 2023-01-01 00:00:00 MST        0
2 USGS       09482495 2023-01-01 00:15:00 MST        0
3 USGS       09482495 2023-01-01 00:30:00 MST        0
4 USGS       09482495 2023-01-01 00:45:00 MST        0
5 USGS       09482495 2023-01-01 01:00:00 MST        0
6 USGS       09482495 2023-01-01 01:15:00 MST        0
7 USGS       09482495 2023-01-01 01:30:00 MST        0
8 USGS       09482495 2023-01-01 01:45:00 MST        0
9 USGS       09482495 2023-01-01 02:00:00 MST        0
10 USGS      09482495 2023-01-01 02:15:00 MST        0
# i 34,693 more rows
# i 2 more variables: qualification_code <chr>, Date <date>
```

Let's create columns for each segment.

```
water <- water %>%
  mutate(Year = year(datetime),
        Month = month(datetime, label = TRUE),
        Day = day(datetime),
        Hour = hour(datetime),
        Minutes = minute(datetime),
        Seconds = second(datetime))

water

# A tibble: 34,703 x 12
  agency_cd site_no   datetime           tz_cd discharge_ft3_s
  <chr>      <chr>    <dttm>          <chr>      <dbl>
```

```

1 USGS      09482495 2023-01-01 00:00:00 MST      0
2 USGS      09482495 2023-01-01 00:15:00 MST      0
3 USGS      09482495 2023-01-01 00:30:00 MST      0
4 USGS      09482495 2023-01-01 00:45:00 MST      0
5 USGS      09482495 2023-01-01 01:00:00 MST      0
6 USGS      09482495 2023-01-01 01:15:00 MST      0
7 USGS      09482495 2023-01-01 01:30:00 MST      0
8 USGS      09482495 2023-01-01 01:45:00 MST      0
9 USGS      09482495 2023-01-01 02:00:00 MST      0
10 USGS     09482495 2023-01-01 02:15:00 MST      0
# i 34,693 more rows
# i 7 more variables: qualification_code <chr>, Year <dbl>, Month <ord>,
#   Day <int>, Hour <int>, Minutes <int>, Seconds <dbl>

```

Now, we can calculate the average value per month.

```

water %>%
  group_by(Month) %>%
  summarise(mean = mean(discharge_ft3_s))

# A tibble: 12 x 2
  Month    mean
  <ord>    <dbl>
1 Jan     0.220
2 Feb     1.46
3 Mar     3.27
4 Apr     3.40
5 May     0.0442
6 Jun     0
7 Jul     0.0222
8 Aug     0.256
9 Sep     0
10 Oct    0
11 Nov    0
12 Dec    0.0886

```

We can also create new columns such as day of year (DOY). This is the day of year, running from 0-365. One common use for DOY is in phenology, measuring the annual timing of life history stages. The function to calculate DOY is `yday()`.

```

water %>%
  mutate(DOY = yday(datetime), .after = tz_cd)

# A tibble: 34,703 x 13
  agency_cd site_no   datetime           tz_cd  DOY discharge_ft3_s
  <chr>     <chr>    <dttm>          <chr> <dbl>            <dbl>
1 USGS      09482495 2023-01-01 00:00:00 MST    1        0
2 USGS      09482495 2023-01-01 00:15:00 MST    1        0
3 USGS      09482495 2023-01-01 00:30:00 MST    1        0
4 USGS      09482495 2023-01-01 00:45:00 MST    1        0
5 USGS      09482495 2023-01-01 01:00:00 MST    1        0
6 USGS      09482495 2023-01-01 01:15:00 MST    1        0
7 USGS      09482495 2023-01-01 01:30:00 MST    1        0
8 USGS      09482495 2023-01-01 01:45:00 MST    1        0
9 USGS      09482495 2023-01-01 02:00:00 MST    1        0
10 USGS     09482495 2023-01-01 02:15:00 MST   1        0
# i 34,693 more rows
# i 7 more variables: qualification_code <chr>, Year <dbl>, Month <ord>,
#   Day <int>, Hour <int>, Minutes <int>, Seconds <dbl>

```

Creating Dates from Segments

What if you are given a dataset that has separate year, month, and day columns that you need to combine? If you recall, this is how the data frame for the Portal rodents is set up.

```

surveys <- read_csv("surveys.csv")

Rows: 35549 Columns: 9
-- Column specification -----
Delimiter: ","
chr (2): species_id, sex
dbl (7): record_id, mo, day, yr, plot_id, hindfoot_length, weight

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```

```

surveys
```

```

# A tibble: 35,549 x 9
  record_id    mo    day    yr plot_id species_id sex  hindfoot_length weight
  <dbl>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
```

```

<dbl> <dbl> <dbl> <dbl> <dbl> <chr> <chr> <dbl> <dbl>
1      1     7    16  1977      2  NL   M       32  NA
2      2     7    16  1977      3  NL   M       33  NA
3      3     7    16  1977      2  DM   F       37  NA
4      4     7    16  1977      7  DM   M       36  NA
5      5     7    16  1977      3  DM   M       35  NA
6      6     7    16  1977      1  PF   M       14  NA
7      7     7    16  1977      2  PE   F       NA  NA
8      8     7    16  1977      1  DM   M       37  NA
9      9     7    16  1977      1  DM   F       34  NA
10     10    7    16  1977      6  PF   F       20  NA
# i 35,539 more rows

```

We can bring these columns together to form a date column using the `make_date()` function.

```

surveys <- surveys %>%
  mutate(date = make_date(year = yr,
                         month = mo,
                         day = day))
surveys

# A tibble: 35,549 x 10
  record_id   mo   day   yr plot_id species_id sex hindfoot_length weight
  <dbl> <dbl> <dbl> <dbl> <dbl> <chr>   <chr>        <dbl> <dbl>
1      1     7    16  1977      2  NL   M       32  NA
2      2     7    16  1977      3  NL   M       33  NA
3      3     7    16  1977      2  DM   F       37  NA
4      4     7    16  1977      7  DM   M       36  NA
5      5     7    16  1977      3  DM   M       35  NA
6      6     7    16  1977      1  PF   M       14  NA
7      7     7    16  1977      2  PE   F       NA  NA
8      8     7    16  1977      1  DM   M       37  NA
9      9     7    16  1977      1  DM   F       34  NA
10     10    7    16  1977      6  PF   F       20  NA
# i 35,539 more rows
# i 1 more variable: date <date>

```

Performing Calculations

Because of the unique way that dates and datetimes are stored in R, we can also perform calculations with them.

In particular, `lubridate` has functions to calculate intervals between dates and datetimes.

```
start <- min(surveys$date, na.rm = TRUE)
end <- max(surveys$date, na.rm = TRUE)
```

```
date_interval <- interval(start, end)
date_interval
```

```
[1] 1977-07-16 UTC--2002-12-31 UTC
```

```
as.duration(date_interval)
```

```
[1] "803433600s (~25.46 years)"
```

This lesson has really only scraped the surface of working with dates, but it should get you started if you need to do so in the future!